

FC Layer Mapping - Detailed Explanation

1. Gated Attention Mechanism (Controlled Influence of Each Encoder)

Instead of directly concatenating **E_lab** (Lab Report Embedding) and **E_conv** (Conversation Embedding), introduce a **gating mechanism**:

Step 1: Compute Attention Scores for Each Encoder

- Define learnable **gating weights** for **Lab vs. Conversation**:
- - $\alpha_{\text{lab}} = \sigma(W_{\text{lab}} \cdot E_{\text{lab}} + b_{\text{lab}})$
 - $\alpha_{\text{conv}} = \sigma(W_{\text{conv}} \cdot E_{\text{conv}} + b_{\text{conv}})$
- σ is a **sigmoid function** (keeps values between 0 and 1).
- $W_{\text{lab}}, W_{\text{conv}}$ are **learnable weights**.
- $\alpha_{\text{lab}}, \alpha_{\text{conv}}$ represent how much each encoder contributes.

Step 2: Weighted Fusion of Embeddings

$$E_{\text{fused}} = \alpha_{\text{lab}} \cdot E_{\text{lab}} + \alpha_{\text{conv}} \cdot E_{\text{conv}}$$

♦ Impact:

- If the lab report is **more reliable**, the model **learns to assign a higher weight to E_lab**
- If conversation symptoms provide **critical missing details**, the model adjusts accordingly.
- If **conflicting information** exists (lab suggests Malaria, but symptoms don't fully match), the model **reduces attention on conversation embeddings** to avoid misclassification.

2. Feature Attribution (Understanding Which Encoder Contributed More to a Prediction)

Once the model has made a prediction, you can analyze **how much each encoder influenced the final decision**.

Method 1: Gradient-Based Attribution

Compute gradients of the output probability $p_{j|p}$ (for a disease j) w.r.t. each embedding:

$$\partial p_j / \partial E_{\text{lab}}, \partial p_j / \partial E_{\text{conv}}$$

Interpretation:

- If $\partial p_j / \partial E_{lab}$ is high, the neuron learned mostly from lab reports.
- If $\partial p_j / \partial E_{conv}$ is high, the neuron relied more on conversation symptoms.

Method 2 : Attention Weights Analysis (If Using Attention Mechanism)

♦ How it Works:

- If you use **self-attention** or **cross-attention** in your model, you can extract the **attention scores** for each encoder.
- The **higher the attention score for Lab vs. Conversation**, the more influence that encoder had.

♦ How to Implement:

Extract the attention weight matrix after fusion:

$\alpha_{lab}, \alpha_{conv} = \text{softmax}(W_{att} \cdot [E_{lab}, E_{conv}])$

- Higher α_{lab} → More reliance on lab reports.
- Higher α_{conv} → More reliance on symptoms.
-

♦ Impact:

- ✓ Gives a **direct interpretable score** of **encoder contribution** for each prediction.
 - ✓ More reliable than gradient-based methods like **Saliency Maps**, which can be noisy.
-

Method 3 : Cosine Similarity Between Neuron Weights and Encoders

♦ How it Works:

- Measures **how aligned** each neuron's learned weights are with **Lab vs. Conversation embeddings**.

♦ How to Implement:

- Compute **cosine similarity** between **ICD neuron weights** and each encoder embedding:

Similarity calculations:

- **Similarity_lab** = $(W_j \cdot E_{lab}) / (\|W_j\| \|E_{lab}\|)$
- **Similarity_conv** = $(W_j \cdot E_{conv}) / (\|W_j\| \|E_{conv}\|)$

- Higher similarity → The **ICD neuron relies more on that encoder**.

♦ **Impact:**

- ✓ Helps quantify **which encoder's features contribute most** to the disease prediction.
 - ✓ **Easy to implement**, works even without attention layers.
-

3. Weight-Based Analysis (Neuron Activation Per Encoder)

- Analyze how strongly each neuron **activates** in response to **E_lab vs. E_conv**.
- Compute **cosine similarity** between **final neuron weights** and each encoder embedding:

Here is the properly formatted version in plain text:

$$\text{Similarity_lab} = (\mathbf{W_j} \cdot \mathbf{E_{lab}}) / (\|\mathbf{W_j}\| \|\mathbf{E_{lab}}\|)$$

$$\text{Similarity_conv} = (\mathbf{W_j} \cdot \mathbf{E_{conv}}) / (\|\mathbf{W_j}\| \|\mathbf{E_{conv}}\|)$$

♦ **Interpretation:**

- If **Similarity_lab > Similarity_conv**, neuron relies **more on lab reports**.
- If **Similarity_conv > Similarity_lab**, neuron relies **more on symptoms**.

You can **tune the training loss** to **increase weight on lab reports** when needed.


4. Preventing Vague Predictions (Thresholding & Confidence Calibration)

- **Use Probability Thresholding:** Ensure that **diseases are predicted only if confidence is above a threshold** (e.g., 0.7).
- **Calibration:** Train the model to **output uncertainty scores** to detect **low-confidence cases**.

♦ **Impact:**

- If conversation symptoms **contradict the lab report**, the model should **decrease prediction confidence** and **raise a flag for doctor review** instead of making a vague prediction.

How the Conventional Method Works (Directly Sending Codes to Loss Function for Multilabel Classification)



This approach **directly maps ICD codes to the final layer neurons** and optimizes prediction using **multilabel classification**.

1. Model Structure

- **Encoders**
 - **Lab Report Encoder (E_{lab} , 768-dim)**
 - **Conversation Encoder (E_{conv} , 1024-dim)**
 - **Fused Embedding: E_{fused}** (concatenated or dimensionally reduced to 1536-dim or 1792-dim)
- **Fully Connected (FC) Layer**
 - **Number of neurons = Number of ICD codes** (e.g., 2000 neurons for 2000 diseases)
 - **Each neuron corresponds to one disease.**
 - **Output = Logits for each ICD code**

2. Forward Pass (How the Model Predicts)

Computing Logits

- $z_j = W_j \cdot E_{fused} + b_j$
 - $W_j \rightarrow$ Weights for disease j
 - $E_{fused} \rightarrow$ Input embedding
 - $b_j \rightarrow$ Bias term
 - $z_j \rightarrow$ Logit (raw prediction score)

Applying Activation Function (Sigmoid for Multilabel Classification)

- $p_j = 1 / (1 + e^{(-z_j)})$

Each neuron independently predicts a probability for its corresponding disease.

3. Backpropagation (How the Model Learns)

Loss Function: Binary Cross-Entropy (BCE)

$$L = - \sum_{j=1}^C [y_j \log(p_j) + (1 - y_j) \log(1 - p_j)]$$

- $C \rightarrow$ Total number of ICD codes
- $y_j = 1$ if the disease is present, otherwise 0

- **Weight Updates:**
If the model predicts Malaria but the patient actually has Typhoid, it reduces W_{malaria} and increases W_{typhoid} .
 - With each training iteration, the predictions are refined.
-

4. Why Does This Work for Fast Deployment?

- ◆ **Simple to implement** (Standard classification model with sigmoid + BCE loss).
- ◆ **Trains quickly** as it doesn't require attention mechanisms.
- ◆ **Best for deployment**
- ◆ **Control over Conversation and Lab Reports Embedding**
- ◆ **Could alter their influence on the model easily**
- ◆ **Multi-disease support** (predicts multiple ICD codes independently).